Intergenerational Ethics

Tim Mulgan

Unless something goes drastically wrong in the next few centuries, most of the people who will ever live are yet to be born. Our actions have little impact on those who are dead, considerable impact on those currently alive, and potentially enormous impact on those who will live in the future. Our decisions affect who those future people will be, and even if there will be any future people at all. The threat of environmental crisis gives us some inkling of the magnitude of our impact on future generations. Only in the last few decades have moral philosophers really begun to grapple with the complexities of intergenerational ethics. Underlying their often technical debates are some of the deepest moral questions. What makes life worth living? What do we owe to our descendants? How do we balance their needs against our own?

Philosophical discussion of intergenerational ethics focuses on ways that our relations with future people differ from relations between contemporaries. There is only a distinct subject if there is a distinctive subject matter. One difference is that, while everyday decisions impact on other people, only in intergenerational ethics do our decisions affect who will begin to exist – and how many people will begin to exist. Different population or family planning policies bring different sets of people – and different numbers of people – into existence.

A related feature of intergenerational ethics is a stark lack of reciprocity. While our decisions affect the lives of future people, their actions have no impact on us. We can do a great deal for (or to) posterity but posterity cannot do anything for (or to) us. If we think of morality as a bargain or contract, then it seems we have no obligations to future people at all.

These distinctive features of intergenerational ethics raise three central questions: (1) Do we have any obligations to future people? (2) If so, what grounds those obligations? (3) Finally, what obligations do we have? While the third question is clearly the most urgent from a practical point of view, most philosophical attention has focused on the first two questions. This essay examines the two dominant traditions: utilitarianism and social contract. It first asks why intergenerational ethics is such a recent subject.

Why Did Philosophers Ignore the Future?

Until very recently, moral philosophy concentrated on interactions between contemporaries. Future generations were only ever an afterthought. To see why, we first distinguish three general sets of background assumptions one might bring to the study of intergenerational ethics.

The International Encyclopedia of Ethics. Edited by Hugh LaFollette, print pages 2680–2693.

DOI: 10.1002/ 9781444367072.wbiee088

The Optimistic Model. We will certainly leave future generations better-off than ourselves, and we ought to do so. The central question is how much better-off we should leave them, and at what sacrifice to ourselves.

The Stasis Model. We can leave future generations at least as well-off as ourselves, at some cost to ourselves. We ought to do so. However, we cannot make future generations better-off than ourselves.

The Pessimistic Model. We cannot avoid leaving future generations worse-off than ourselves. By making significant sacrifices, we can reduce the extent to which they are worse-off than us. The crucial questions are: How much worse-off are we allowed to leave them, and how much are we required to sacrifice?

Until recently, the optimistic model was largely taken for granted by moral and political philosophers. This optimism enables contemporary political theory to ignore the future. We need only look after ourselves, do what is best for present people, and then bequeath our stable liberal democratic institutions, thriving economy, and scientific advances to future people. What is good for us is also good for them. There is no conflict between present and future. The classic example is John Rawls (*see* RAWLS, JOHN). Rawls (1971) devotes just 10 pages to justice between generations; and his only intergenerational ethical question is the just savings problem: How much better-off should we leave our descendants?

After the oil shocks of the early 1970s, the realization that the standard of living of the developed nations currently requires the rapid depletion of nonrenewable resources lead philosophers to explore the implications of the stasis model. Barry (1978) explored various constraints on the appropriation and consumption of resources modeled on John Locke's principle that those who appropriate resources must leave "as much and as good" for others (*see* LIBERTARIANISM; LOCKE, JOHN). Even if we cannot ensure that future generations are better-off, we should at least not leave them worse-off – perhaps by compensating them with the provision of superior technology.

The stasis model raises many practical issues. But it may not seem very theoretically troubling. If we can leave future people as well-off as ourselves, then obviously we should do so. However, the stasis model is important, because it raises the possibility of a conflict between generations. Once we admit that possibility, we need to know how to think about such conflicts. And we need to talk in terms of obligations to future people.

One exception to the general optimism of political philosophy has been the recognition, since Robert Malthus in the early nineteenth century, of the threat that overpopulation will reduce average welfare. Overpopulation has also featured prominently in philosophical discussion of famine in the developing world, especially in Africa, since the 1960s. In recent years, debates over carbon emissions, greenhouse gases, ozone depletion, and the serious threat of climate change have all brought the pessimistic model to the fore in relation to the developed world (*see* ENVIRONMENTAL ETHICS). It may already be too late to prevent future generations being worse-off

than us. Furthermore, any feasible proposal to minimize the harmful effects of climate change will involve a major reduction in the standard of living of affluent people in developed countries. The pessimistic model raises many difficult and novel ethical and political questions – and brings a new urgency to the search for foundations and principles in intergenerational ethics.

Utilitarian Intergenerational Ethics

The recent philosophical literature on intergenerational ethics is dominated by utilitarianism (*see* UTILITARIANISM). Unlike some competitors, utilitarianism has a natural and compelling explanation of the existence of obligations to future people. Morality is based on the promotion of human well-being (*see* HAPPINESS; WELL-BEING). Human lives matter equally wherever – and whenever – they are lived. If there will be future people, and if our actions can impact on their quality of life, then we have obligations with respect to them. Those obligations are no different to those we have regarding present people.

More controversially, utilitarianism also explains why we have an obligation to ensure that there are future people. If human well-being is a good thing, then it is good that there be happy people – and not just that any people who happen to exist are happy. Some opponents regard this implication as a defect of utilitarianism. They argue that the value of happiness is conditional on the existence of people. If people exist, then it is good that they be happy – but it is not good that there be people (Narveson 1967).

Utilitarianism seems, if anything, to give us too many obligations. Given the enormous numbers of future people, and the significance of our potential impact on them, it looks as if intergenerational ethics will swamp all other ethical considerations. Intergenerational ethics is thus a very striking case of a perennial problem for utilitarianism – its demandingness. Consider the analogous case of people who already exist and live in poverty in distant lands. There are very many such people, and I am in a position to assist many of them. But then my utilitarian obligations to them will swamp all my personal projects, and my special obligations to my nearest and dearest. Some utilitarians respond by trying to develop a moderate utilitarian intergenerational ethic – taking existing solutions to the demandingness problem and applying them to future people (Mulgan 2006).

The primary focus of recent utilitarian debate, however, concerns a prior question: the issue of aggregation. We can introduce this issue via a pair of distinctions introduced in Parfit (1984: 351–441) – generally regarded as the starting point for all recent intergenerational ethics. Parfit first distinguishes two kinds of moral choice: *same people* (where our actions affect what will happen to people in the future, but not which people will come to exist) and *different people* (where our actions do affect who will come to exist in the future).

Utilitarianism treats same people and different people choices identically. What matters is how happy people are, not who they are. In Parfit's terminology, utilitarians endorse a *no difference view* (1984: 367). If A and B are two situations, and if the

only difference between them is that A is a different people choice and B is a same people choice, then there is no moral difference whatsoever between A and B. Under the no difference view, different people choices do not, *per se*, present any new ethical issues. Intergenerational ethics looks easy. We just take our familiar principles and apply them to new cases. Unfortunately, things are not so simple. As Parfit notes, we can further divide different people choices into *same number* (where our choice affects who exists, but not how many people exist) and *different number* (where we decide how many people ever exist). Different number choices raise many new difficulties. These centre on a seemingly very abstract question. Suppose you could create any possible world, with any possible population. Which world should you choose? Because they base morality on the maximization of happiness, utilitarians obviously need to answer this question (*see* GREATEST HAPPINESS PRINCIPLE). If we aim to produce the best amount of happiness, then we must decide what counts as the best happiness. Utilitarians need a theory of aggregation – taking us from the values of individual lives to the value of a population as a whole (*see* POPULATION).

As far as possible, when constructing a theory of aggregation, utilitarians seek to remain neutral as to the nature of happiness. In these discussions, "happiness" is often used as a generic place-holder – it refers to whatever makes life worth living. Utilitarian aggregation is exclusively about human well-being. This may seem unacceptably parochial, speciesist, or anthropocentric (*see* ANTHROPOCENTRISM). What about ecological values or the welfare of animals? However, utilitarian values can be combined with nonhuman values. If we believe that human well-being is one value, then we still need an account of aggregation – we can then supplement it by adding the aggregate value of animal lives, or else we can combine aggregated happiness with other ecological values. The question of how to balance these different values is a key issue in environmental ethics.

The utilitarian tradition offers two contrasting ways to aggregate the values of human lives. On the total view, the best outcome contains the greatest total amount of happiness. On the average view, the best outcome contains the highest average level of happiness. The two are not always clearly distinguished. This is understandable. In any same-people choice – or any choice where numbers are not at stake – the two views must coincide. Whatever maximizes the total also maximizes the average. But the two views can come apart in a different-number choice. Suppose we must decide between two population policies for our society. One will lead to a large population with a moderate average level of happiness, while the other would produce a smaller population with a higher average level of happiness. We calculate that the former policy offers greater total happiness, and the latter offers higher average happiness. If we confine our interest to human happiness, which outcome is better?

The total view is the simplest theory of aggregation. It has also been the most popular among utilitarian philosophers (economists, by contrast, often favor the average view). The basic argument for the total view is simple. If we value happiness, then presumably we should aim to produce as much happiness as possible. The most

famous problem for the total view is an inference that goes back to Henry Sidgwick (*see* SIDGWICK, HENRY). It takes its modern name from Parfit (1984: 388).

The repugnant conclusion (*see* REPUGNANT CONCLUSION). For any possible population of at least 10 billion people, all with a very high quality of life, there must be some much larger imaginable population whose existence, if other things are equal, would be better, even though its members have lives that are barely worth living.

To see why the total view implies the repugnant conclusion, begin with a world where 10 billion people all have extremely good lives. Call it A. Imagine a second world, with twice as many people, each of whom is more than half as happy as the people in A. Call this new world B. Total happiness in B exceeds that in A. Now repeat this process until we reach a world with a vast population whose lives are barely worth living. Call this world Z. As each step increases total happiness, Z must be better than A.

Parfit (1984: 390) finds this conclusion "intrinsically repugnant." If the repugnant conclusion is a consequence of the total view, then the total view is unacceptable. The repugnant conclusion is one of the organizing problems of contemporary intergenerational ethics – most philosophers begin their discussions by saying how they will deal with it. They either reject Parfit's intuition that A is better than Z or they reject the total view. The repugnant conclusion is a classic example where a thought experiment is alleged to provide a decisive counterexample to a philosophical view. It is worth exploring both in its own right, and also to illustrate the main approaches in contemporary intergenerational ethics.

Some philosophers reject intuitions altogether. What does it matter if a conclusion "appears" repugnant, so long as it follows from well-established premises? These philosophers then face the challenge of explaining what could ground ethical premises other than a moral intuition of some kind. Most often, when philosophers say that they reject intuitions, this means that they reject some intuitions in favor of others. Nonutilitarians have the option of rejecting all intuitions about the comparative value of possible futures. They have no need to decide whether A is better or worse than Z. However, utilitarians do not have this option. Without a theory of aggregation, their moral theory lacks foundation. A more modest view – one that is open to utilitarians – is to reject all intuitions regarding very large numbers. For instance, Broome (2004: 57–8) argues that "we have no reason to trust anyone's intuitions about very large numbers, however excellent their philosophy. Even the best philosophers cannot get an intuitive grasp of, say, tens of billions of people." We need not abandon intergenerational ethics – nor abandon moral intuitions altogether. We should rely instead on a theory built on our everyday intuitions. Broome argues that the best such theory is the total view.

A less radical response to any alleged counterexample is simply to reject the specific intuition underlying it. In the case of the repugnant conclusion, Ng (1989) objects that we privilege our own perspective and are guilty of "misplaced partiality." We picture the A-lives as similar to our own, and imagine the A-people choosing between A and Z. If we were more impartial, we might see that Z contains more total value than A, and is thus better.

Another general response is to elaborate the thought experiment – highlighting suppressed background assumptions that color our intuitions. Many proponents of the total view defend the repugnant conclusion by examining the Z-world more closely. By definition, the Z-lives are "barely worth living." So we need to know what such lives are like. On the total view, we should create an extra life whenever doing so would raise the total happiness – whenever the extra life itself is worth living. If we imagine a numerical scale of well-being, then the lives in Z must be above zero.

Think about what it means to say that a life is "barely worth living," as the Z-lives are meant to be. This phrase can evocate a life of frustration and pain – one that we would rather not live at all. But, if the Z-lives are like that, then the total view does not conclude that Z is better than A. Parfit (1986b: 148) describes the Z-lives as consisting of nothing but muzak and potatoes. If they are human lives, then it is natural to suppose that such lives also contain negative elements – such as boredom, frustration, or lack of accomplishment and friendship. These features reduce the value of a life. A friendless underachieving human is badly off in a way that a friendless slug is not. We may well feel that a muzak and potatoes life is well below zero.

Utilitarianism says that we should create future people only if (future) aggregate human happiness is positive. If the zero level is so high that most people fall below it – or if the minority who fall below zero have very miserable lives – then it might be better if there were no future people. Utilitarianism would then tell us to bring the human story to a close.

Some utilitarians do defend the total view. But others agree with Parfit's intuition, and seek alternatives. The simplest is the average view. This easily avoids the repugnant conclusion, as A has a higher average happiness than Z. If we average over everyone alive in the future, then the average view tells us to kill anyone whose happiness is below average. To avoid this repellent consequence, we average over all those who will ever live. Killing someone typically makes their life go worse; it does not make it the case that they never existed.

The average view faces other objections that are harder to dissolve. Many of these are variants of the *hermit problem*. Suppose everyone in the cosmos is extremely happy. On a distant uninhabited planet, we create a new person. His life, while very good, is slightly below the cosmic average. The average view says that we have made things worse; and that what we ought to do depends on the happiness of people in distant corners of the cosmos, with whom our hermit will never interact – as the value of those distant lives affects the cosmic average. Both claims seem intuitively implausible. As Parfit (1984: 420) puts it, the mere addition of lives worth living cannot make things worse; and our moral decisions should not depend on how happy the ancient Egyptians were.

The hermit problem plays a similar dialectical role to the repugnant conclusion. Defenders of the average view have the same broad options. They can reject the intuition or deny that this result follows from their theory. One popular response is to limit our calculation of the average happiness to those affected by our actions – thus

removing the need to take account of the welfare of people in the distant past or on distant planets.

Another popular account of aggregation is the lexical view (*see* INCOMMENSURABILITY [AND INCOMPARABILITY]). Suppose you enjoy both Mozart and muzak. Someone offers you a choice between one day of Mozart and as much muzak as you like. You opt for the former, because no amount of muzak could match the smallest amount of Mozart. Philosophers would say you believe that Mozart is lexically superior to muzak.

Lexicality avoids the repugnant conclusion. Suppose the creatures in A and Z belong to different species. Perhaps A contains flourishing human beings while Z is full of slugs. A is better, because 10 billion human lives are more valuable than any number of slug lives. To return to our original comparison, a lexical view of human well-being would hold that 10 billion flourishing human lives trump any number of human lives that are barely worth living. A is better than Z.

The most worrying problem for any lexical account is Parfit's *continuum objection*:

> Mozart and Muzak … seem to be in quite different categories. But there is a fairly smooth continuum between these two. Though Haydn is not as good as Mozart, he is very good. And there is other music which is not far below Haydn's, other music not far below this, and so on. Similar claims apply to the … other things which give most to the value of life. … Since this is so, it may be hard to defend the view that what is best has more value than any amount of what is nearly as good. (Parfit 1986b: 164)

The challenge for the lexical view is to tell us where to draw the line – and why.

The philosophical literature contains many other theories of aggregation. However, these all face similar problems to the three theories we have discussed. One focus of debate is Parfit's *mere addition paradox*, which shows that we cannot avoid the repugnant conclusion and at the same time claim that the mere addition of happy lives never makes things worse. Temkin (1987) concludes that it is impossible to construct an intuitively plausible theory, because our intuitions themselves are inconsistent (*see* INTRANSITIVITY). Dasgupta (1994) and Roberts (2002) propose relativized models of value, where we evaluate different possible worlds relative to the interests of the people who live in them. Perhaps A is better than B from one perspective, while B is better from another.

Intergenerational Social Contracts

One influential strand of Western political philosophy pictures both morality and justice as a contract between rational individuals (*see* SOCIAL CONTRACT). But how can we bargain with future people when their very existence is in our hands? If morality depends on contract, then surely there can be no intergenerational ethics. Some are prepared to bite the bullet. We may choose to take future people into account. And, if some of our contemporaries happen to care about future people (perhaps their own

distant descendants), then we owe it to those contemporaries to consider the interests of those future people. But we have no obligations to future people (Heyd 1992). Consider a time bomb that devastates people in the distant future but has no direct impact until then. Suppose the people who will be affected are so far in the future that no one alive today cares for them. Is it wrong to plant a time bomb? If so, is this as wrong as planting a bomb that would cause the same devastation today?

Here we seem to reach intuitive bedrock. Utilitarians advocate strict temporal neutrality. Planting a time bomb is just as wrong as planting a bomb that explodes today. Others may feel that planting a time bomb violates no obligations – even if they would not plant one themselves. Others fall somewhere in between. Planting a time bomb is wrong – but not as wrong as planting a present bomb. Future people matter, but not as much as present people.

This intermediate view is captured in economics by the social discount rate – where we discount costs or benefits that will occur in the future (*see* ECONOMICS AND ETHICS). Social discount rates have enormous impact in the long term. Even a modest discount can mean that it is not economic to spend a few dollars today to avert a global catastrophe in five hundred years' time. To take one striking example: whether cost–benefit analysis concludes that the future benefit of preventing climate change is worth the present cost, depends very largely on our choice of discount rate.

Everyone agrees we should discount for uncertainty – or for such things as the possibility that humanity will be wiped out by an asteroid strike. But some economists go future. They apply a pure time preference – future happiness counts for less simply because it is in the future. The justification is that this reflects how people actually make decisions about the future. We do discount future benefits – both to ourselves and to others. Others disagree. For instance, Cowen and Parfit (1992) argue that a pure time preference has no place in intergenerational ethics.

Suppose we want to construct an intergenerational social contract. We might first note that overlapping generations do interact and bargain. We can then extend our contract indefinitely into the future, using what Gosseries (2001) dubs the *zipper argument*. Suppose we have only three generations: G1, G2, and G3. G1 and G3 do not interact, but G2 interacts with both G1 and G3. G2 know that they will have to bargain with G3. So G2 will take G3's interests into account when bargaining with G1 – and ask G1 not to leave a bomb that will devastate G3.

This ingenious argument is problematic. Standard social contracts assume self-interest. But then why will G2 object to a time bomb that will impact only on G3? On the contrary, G2 might welcome the bomb, as it strengthens G2's position against G3 ("If you don't give us what we want, we will not defuse the time bomb"). G2 would then ask G1 to plant such a bomb – and so time bombs would be morally desirable (Mulgan 2006: 28–32).

Many philosophers conclude that a self-interested contract is a poor foundation for morality. They prefer hypothetical or idealized contracts. The classic recent example is Rawls (1971), who asks what people would agree to under certain idealized circumstances. Rawls seeks principles of justice everyone can recognize as a fair basis for mutual interaction. These principles are chosen in an original position,

from behind a veil of ignorance. The choosers know what their society will look like if any given principle is adopted, but they do not know who they will be in that society. Imagine a very simple society with two groups: rich and poor. What principles would a rational person choose if they did not know whether they were rich or poor?

In Rawls (1971: 284–93), the parties to the original position belong to the same generation. Unless they care about future people, intergenerational justice will not feature in their principles. Nothing we do to future people – however devastating – could count as unjust. As an egalitarian, this conclusion would be unacceptable for Rawls. He must accommodate future generations. Rawls originally added a motivational assumption. Those in the original position care about their descendants, at least for the next generation or two. This solution is ad hoc. Why allow concern for descendants, when we allow no concern for contemporaries? Furthermore, any realistic motivational assumption only works for a few generations. It thus cannot remove the threat of time bombs. Rawls focuses on savings from one generation to the next, not on longer term environmental or resource issues. This focus was controversial even at the time (1971) – and seems much more problematic now.

Rawls (1993: 273–4) abandoned this solution, and stipulated instead that those in the original position must behave in a way that they would want previous generations to have behaved. Total self-sacrifice is ruled out, as the cost of our sacrifice outweighs the benefits of the sacrifices of others. Total selfishness also fails, as the damage of earlier selfishness outweighs our own freedom to behave as we wish. We need something in between. Unfortunately for Rawls, it is very hard to say what that something will be.

A more intriguing option for Rawls is to extend the veil of ignorance, so that people don't know what generation they belong to (Barry 1989: 179–203). Each generation then cares for the interests of all. But now we must decide who participates in this new original position. Before the present generation decide how they will live, there is no fact of the matter as to who will exist in the future. We thus have two alternatives. The first is to extend the veil of ignorance to include everyone who will ever live. We do know we will live at some point, but we do not know when. This leads to problems analogous to those faced by the total view. The parties will seek to maximize their chances of existing even if the result is Parfit's Z-world. The other alternative includes everyone who might exist. Now we do not even know whether we will ever exist. But this leads to problems similar to those facing the average view, as the parties will opt for a tiny population with the highest possible quality of life.

Like utilitarianism, the social contract seems to need radical revision if it is to ground intergenerational ethics. It is thus no surprise that philosophers seek alternatives to both utilitarianism and social contract. They seek to ground intergenerational ethics in something other than the welfare of individual human beings, or the interactions of contemporaries. Perhaps our obligations are owed not to future people but to past people – who sacrificed so that we could thrive, and would expect us to do the same for those who will come after. Our obligations regarding future people would then actually be debts owed to our ancestors. Or perhaps our moral

obligation is to a community that persists through the generations – or perhaps even to humanity as a whole. One early exploration of the various resources available to communitarian intergenerational ethics is de Shalit (1994; *see* COMMUNITARIANISM). These solutions remove the distinctiveness of intergenerational ethics, by focusing on features it shares with relations between contemporaries.

Person-Affecting Principles

The search for foundations for intergenerational ethics is problematic. An alternative is to look instead for concrete principles of intergenerational ethics. One place to start is with a set of principles rejected by Parfit (1984: 351–79). Parfit claims that different people choices are more frequent than we think; and that traditional moral theories are designed for same-people choices, and must be amended for different-people choices. These claims constitute the nonidentity problem, so called because those who exist in one possible future are not (numerically) identical to those who exist in another. To illustrate the nonidentity problem, consider an example adapted from Parfit (1984: 371). Suppose we must choose an energy policy. Should we bury nuclear waste in a desert, or opt for a safer alternative? Different policies produce different patterns of migration. Suppose we choose the riskier option. It leads to a catastrophe in several centuries' time. Intuitively, we have done something wrong. But now take any particular individual killed by that catastrophe. It is almost certain that she herself would never have existed if we had chosen the other policy. So no particular person is worse-off than she herself would have been if we had chosen differently.

Parfit (1984: 362) also offers a less dramatic example, where we choose between depleting and conserving natural resources. Those who live in the future are worse-off under depletion, but they still have lives worth living. Because they would not have existed under a policy of conservation, they seem to have no complaint.

Nonidentity is thus a significant problem for any person-affecting principle – any principle that says an action can only be wrong if some particular person is worse off than they would otherwise have been. In a different people choice, whatever we do, no particular individual is worse-off than she would otherwise have been – as she would otherwise not have existed, and we cannot compare existence with nonexistence. No person-affecting theory can ever condemn any different people choice, however horrific the resulting lives. Even if we create a person whose life contains nothing but excruciating agony, we have done nothing wrong.

The nonidentity problem has implications for many areas of practical ethics. In medical ethics, it arises whenever any individual reproductive choice or medical procedure affects the identity of the resulting child. For instance, many people object to new reproductive technologies on the grounds that they harm the resulting children. But, if they affect a child's genetic makeup, and if genetic identity is a component of individual identity, then such technologies generate different people choices – and the resulting children would not otherwise have existed. If their lives are worth living, how can they be said to be harmed (*see* GENETIC TESTING;

REPRODUCTIVE TECHNOLOGY; SAVIOR SIBLINGS; WRONGFUL LIFE)? The nonidentity problem also arises in discussions of reparations for historical injustice. Should present people be compensated for some past wrong if they would not have existed in an alternative future where the injustice did not occur? Can the descendents of those who suffered from slavery or colonization consistently complain about injustices without which they themselves would not exist? Or do they need to make their claims on behalf of groups that would otherwise have existed (*see* REPARATIONS)?

The nonidentity problem is often used to defend utilitarianism. Because it adopts a no difference view, utilitarianism is not person-affecting. It can easily say what is wrong with Parfit's risky policy, or with any action that creates less happy people. Others seek to defend the person-affecting view. They first argue that creating a person who experiences nothing but agony is wrong. Such a life is not worth living. And we can reasonably say that such a life is worse for that person than nonexistence – even though there is no sense in which nonexistence would have been better. This enables person-affecting principles to condemn the creation of a person whose life is below zero (McMahan 2001).

Many person-affecting theorists seek stronger obligations. One common defense is as follows. The main intuition underlying any person-affecting theory is that an action is only wrong if someone is wronged. But a person can be wronged even if it is not the case that they would otherwise have been worse-off (Woodward 1986). The classic example is when a person is prevented from boarding a plane because of his race, and the plane goes on to crash. This person has been wronged – even though he would otherwise have died. A person can be wronged by an act leading to her creation, even if her life is worth living and she would otherwise not have existed at all. For instance, if a couple choose to have a disabled child simply to develop their own capacity for compassion, then they wrong that child by treating her as merely a means to their own ends – even if her life is overall worth living. Parfit (1986a) replies that if a future person has a life that is overall worth living, then she will waive her rights – precisely because her only alternative was not to exist at all. Debate then centers on whether rights are the appropriate moral idiom for this situation, and whether such rights can really be waived. Person-affecting theorists are also challenged by cases such as Parfit's depletion versus conservation, where it is much more controversial whether any rights are violated – especially if depletion leaves future people better-off than many present people. (On the other hand, the anti-person-affecting intuition is less compelling in this case. Is depletion really wrong when it leaves future people well-off?)

An alternative approach is to reinterpret person-affecting principles so that they apply directly to different people choices. Kumar (2003) gives the example of principles governing the way a parent should treat her child. What matters here is the relationship between parent and child, not the identity of the child. Suppose a prospective mother decides not to take a drug that would give her a disabled child. Pausing to take the drug would have changed the identity of her

child, so she faced a different person choice. She can still say, "I didn't do that because it would have harmed my child." In this case, the phrase "my child" covers both of the two (numerically distinct) children she might have had. Kumar embeds his broader person-affecting principle in a contractualist framework, thereby offering another possible foundation for intergenerational ethics (*see* CONTRACTUALISM).

Both the foundations and the content of intergenerational ethics are sources of controversy. Intergenerational ethics is a comparatively new area of philosophical inquiry, and is thus likely to see significant developments in the future.

***See also*:** ANTHROPOCENTRISM; COMMUNITARIANISM; CONTRACTUALISM; ECONOMICS AND ETHICS; ENVIRONMENTAL ETHICS; GENETIC TESTING; GREATEST HAPPINESS PRINCIPLE; HAPPINESS; INCOMMENSURABILITY (AND INCOMPARABILITY); INTRANSITIVITY; LIBERTARIANISM; LOCKE, JOHN; POPULATION; RAWLS, JOHN; REPARATIONS; REPRODUCTIVE TECHNOLOGY; REPUGNANT CONCLUSION; SAVIOR SIBLINGS; SIDGWICK, HENRY; SOCIAL CONTRACT; UTILITARIANISM; WELL-BEING; WRONGFUL LIFE

REFERENCES

Barry, Brian 1978. "Circumstances of Justice and Future Generations," in R. Sikora and B. Barry (eds.), *Obligations to Future Generations*. Philadelphia: Temple University Press, pp. 204–48.

Barry, Brian 1989. *Theories of Justice*. Berkeley: University of California Press.

Broome, John 2004. *Weighing Lives*. Oxford: Oxford University Press.

Cowen, Tyler, and Derek Parfit 1992. "Against the Social Discount Rate," in Peter Laslett and James Fishkin (eds.), *Justice between Age Groups and Generations*. New Haven: Yale University Press, pp. 144–61.

Dasgupta, Partha 1994. "Savings and Fertility: Ethical Issues," *Philosophy and Public Affairs*, vol. 23, pp. 99–127.

Gosseries, Axel 2001. "What Do We Owe the Next Generation(s)?" *Loyola of Los Angeles Law Review*, vol. 35, pp. 293–354.

Heyd, David 1992. *Genethics: Moral Issues in the Creation of People*. Berkeley: University of California Press.

Kumar, Rahul 2003. "Who Can Be Wronged?" *Philosophy and Public Affairs*, vol. 31, pp. 99–118.

McMahan, Jeff 2001. "Wrongful Life: Paradoxes in the Morality of Causing People to Exist," in John Harris (ed.), *Bioethics*. Oxford: Oxford University Press, pp. 445–75.

Mulgan, Tim 2006. *Future People*. Oxford: Oxford University Press.

Narveson, Jan 1967. "Utilitarianism and New Generations," *Mind*, vol. 76, pp. 62–72.

Ng, Yew-Kwan 1989. "What Should We Do About Future Generations? Impossibility of Parfit's Theory X," *Economics and Philosophy*, vol. 5, pp. 235–53.

Parfit, Derek 1984. *Reasons and Persons*. Oxford: Oxford University Press.

Parfit, Derek 1986a. "Comments," *Ethics*, vol. 96, pp. 832–72.

Parfit, Derek 1986b. "Overpopulation and the Quality of Life," in Peter Singer (ed.), *Applied Ethics*. Oxford: Oxford University Press, pp. 145–64.

Rawls, John 1971. *A Theory of Justice*. Cambridge, MA: Harvard University Press.

Rawls, John 1993. *Political Liberalism*. New York: Columbia University Press.
Roberts, Melinda 2002. "A New Way of Doing the Best We Can: Person-Based Consequentialism and the Equality Problem," *Ethics*, vol. 112, pp. 315–50.
Shalit, Avner de 1994. *Why Posterity Matters: Environmental Policies and Future Generations*. Oxford: Routledge.
Temkin, Larry 1987. "Intransitivity and the Mere Addition Paradox," *Philosophy and Public Affairs*, vol. 16, pp. 138–87.
Woodward, James 1986. "The Non-Identity Problem," *Ethics*, vol. 96, pp. 804–31.

FURTHER READINGS

Arrhenius, Gustaf 1999. "Mutual Advantage Contractarianism and Future Generations," *Theoria*, vol. 1999, pp. 25–35.
Barry, Brian 1977. "Justice between Generations," in P. M. S. Hacker and J. Raz (eds.), *Law, Morality and Society: Essays in Honour of H. L. A. Hart*. Oxford: Oxford University Press, pp. 268–84.
Feinberg, Joel 1986. "Wrongful Life and the Counterfactual Element in Harming," *Social Policy and Philosophy*, vol. 4, pp. 145–78.
Fotion, N., and J. C. Heller (eds.) 1997. *Contingent Future Persons: On the Ethics of Deciding Who Will Live, or Not, in the Future*. Dordrecht: Springer.
Griffin, James 1986. *Well-Being*. Oxford: Oxford University Press.
Hurka, T. 1983. "Value and Population Size," *Ethics*, vol. 93, pp. 496–507.
Kavka, Gregory 1982. "The Paradox of Future Individuals," *Philosophy and Public Affairs*, vol. 11, pp. 93–112.
Meyer, Lukas 2008. "Intergenerational Justice," *Stanford Encyclopedia of Philosophy*. At http://plato.stanford.edu/archives/fall2008/entries/justice-intergenerational/, accessed December 27, 2011.
Roberts, Melinda 1998. *Child Versus Childmaker: Future Persons and Present Duties in Ethics and the Law*. Oxford: Rowman & Littlefield.
Roberts, Melinda, and David Wasserman (eds.) 2009. *Harming Future Persons: Ethics, Genetics and the Nonidentity Problem*. Dordrecht: Springer.
Ryberg, Jesper, and Torbjorn Tannsjo (eds.) 2004. *The Repugnant Conclusion: Essays on Population Ethics*. Dordrecht: Kluwer.
Steinbock, Bonnie 1986. "The Logical Case for 'Wrongful Life'," *Hastings Centre Report*, vol. 16, pp. 15–20.
Vallentyne, Peter 1993. "Utilitarianism and Infinite Utility," *Australasian Journal of Philosophy*, vol. 71, pp. 212–17.